

Reliability and Flow control


TCP revisited

- TCP provides the following abstraction
 - In-order delivery
 - **Reliability**
- Reliability while improving performance requires
 - **Flow control**
 - **Congestion control**

Two common ways for reliability in networking

- Automatic Repeat Request (ARQ)
 - Usually used in all layers.
- Forward Error Correction (FEC)
 - Usually used in lower layers

Simple ARQ solution: Stop and Wait




- Sender doesn't send next packet until she's sure receiver has last packet
- What are the problems with Stop and Wait?

Performance problem with stop and wait

- Problem: “keeping the pipe full”
 - If the bandwidth-delay product [BDP] is much larger than a packet size, the sender will be unable to keep the link busy
 - **Bandwidth-delay product** estimates the amount of data that can be pushed on the pipe

Sliding window



Sliding window

- TCP sends a window size of packets instead of stop-and-wait
- For window size n , sender may transmit n bytes without receiving an ACK
 - After each ACK, the window slides forward

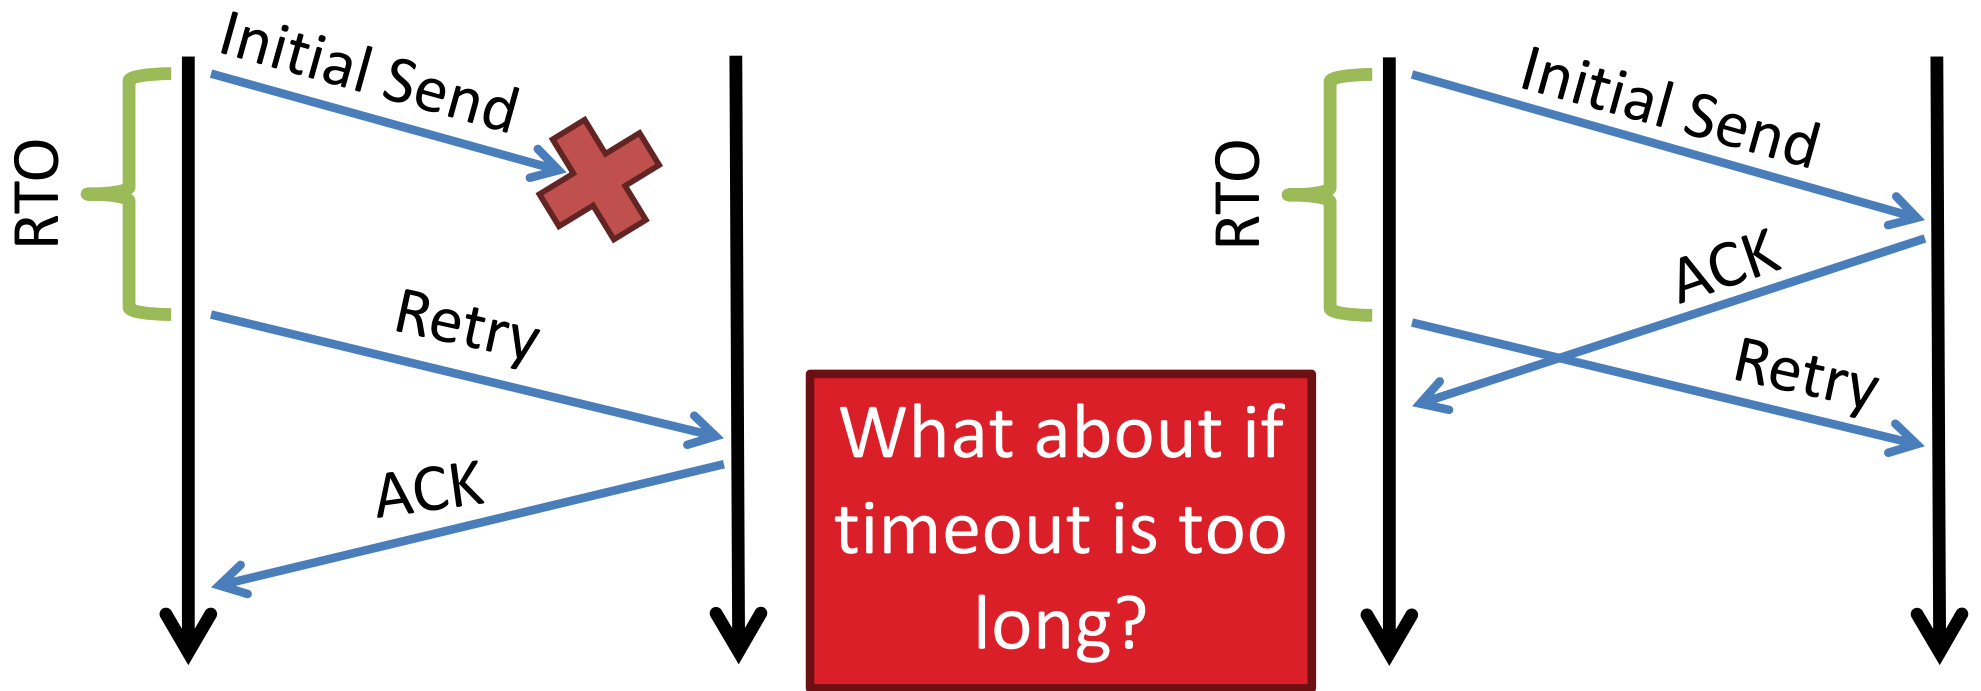
What happens when packets are lost?

Timeouts

- Lost segments detected by sender
 - Use **timeout** to detect missing ACKs
 - What should the timeout depend on?

Retransmission Time Outs (RTO)


- Problem: time-out could be too short



TCP's timeout

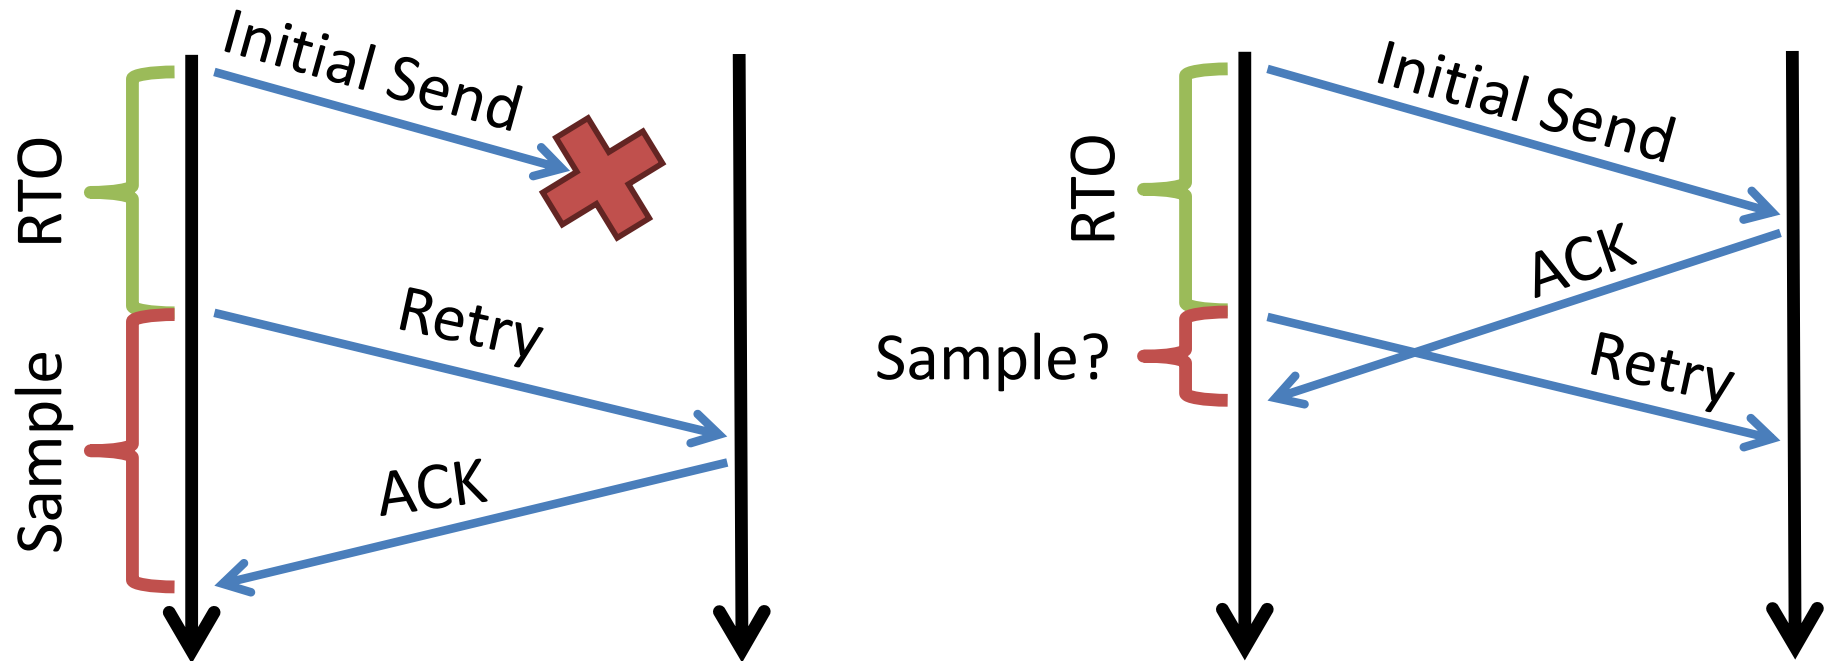
- One idea
 - $RTO = 2 * RTT'$
 - RTT' is the estimated RTT
 - TCP is conservative
- RFC gives the exact RTO measure
 - <https://tools.ietf.org/html/rfc6298>
- How should we estimate RTT?

Round Trip Time Estimation



- $RTT = (1-\alpha) (RTT) + \alpha(\text{new_sample})$
 - Recommended α : 0.125

RTT Sample Ambiguity



- Karn's algorithm: ignore samples for retransmitted segments

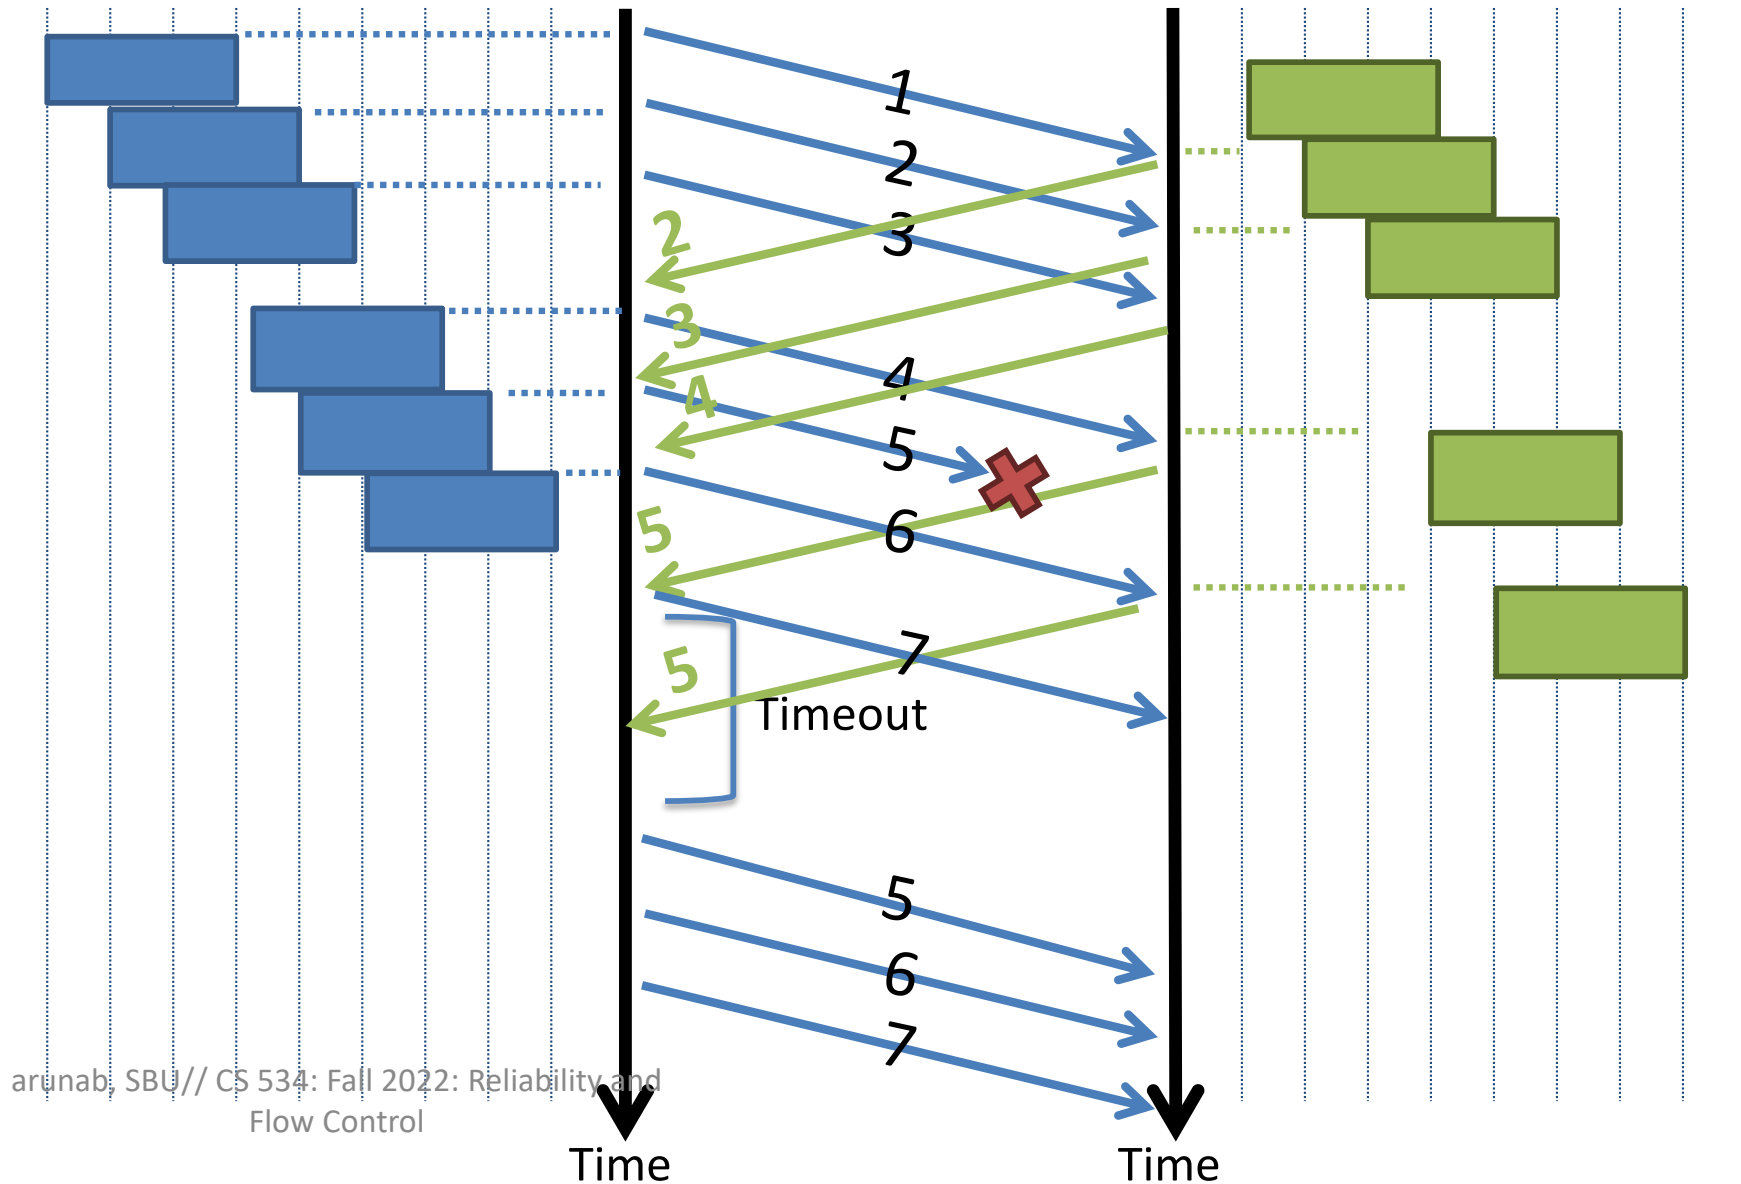
When should the Receiver ACK?

1. ACK *every packet*
2. Use *cumulative ACK*, where an ACK for sequence n implies ACKS for all $k < n$

Later we discuss triple duplicate acks, NACKs, and SACKs

Example: Cumulative acks

Sliding Window = 3



Sliding window problem

- Determine how many packets can be sent into the system
- Send too many
 - Can overwhelm the receiver
 - Can overwhelm the network
- Solution:
 - Flow control (not to overwhelm the receiver)
 - Congestion control (not to overwhelm the network)

Bandwidth delay product (BDP)

- It is theoretically the maximum number of packets you can send before receiving an acknowledgment
- Since acknowledgement takes roughly RTT time to get, the number of packets you can send in one RTT
$$\text{BDP} = \text{bandwidth} \times \text{RTT}$$
- In general, you cannot send BDP packets in the network because there are other bottlenecks
 - Receiver bottleneck and Network bottleneck

Receiver bottleneck and network bottleneck

- Receiver centric: Receiver does not have a big enough buffer to store un-ordered packets
 - Solution: Flow control
- Network centric: The routers do not have capacity to process the packets, causing queue buildup and packet drops
 - Solution: Congestion control

Flow Control

- Problem: how many packets should a sender transmit?
 - Too many packets may overwhelm the receiver
- Solution
 - Receiver tells the sender how big their buffer is
 - This is called the **advertised window**

Flow Control: Sender Side

